a Benchmark construction

253 candidate pairs
new-indication approvals 2023–2026, four agencies

↓ −19 unresolved IDs, or placeholder indication text

234 identifiers resolved
DrugBank ID and a disease ontology ID

↓ −118 answer already an edge in a KG, or drug/disease absent from all three

116 leakage-free and KG-resident
answer is not already an edge; both endpoints present

↓ −17 duplicate pair: the same approval recorded twice

99 unique drug–disease pairs
one query per pair

↓ −19
5 contradicted by the register
4 not an indication change
10 population or severity expansion

80 **register-verified strict set**
every row cited to a regulator record

An independent rebuild from the candidate file, applying the same criteria from scratch, reproduces 71 of these 80. The 9 differences are listed in cleaning_manifest.csv.

b Knowledge-graph coverage of the 80

of 73 pairs resolvable in the coverage audit; 7 combination products excluded

49 covered by all three: the common support on which evidence quality is comparable across graphs
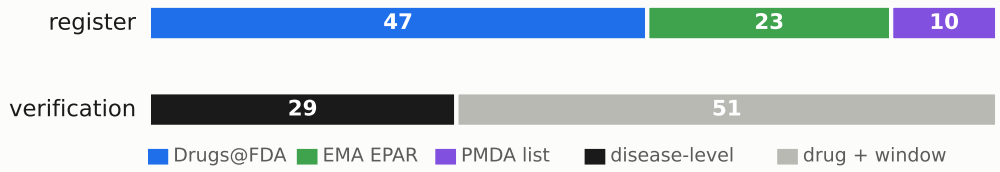
c Evidence behind each pair